Contents

Prerequisiti

Incluso In parte, se abbiamo tempo

Escluso

Preface *page* xvi

Acknowledgments xxiii

Part I Fundamentals and Probability on Events

1 Before We Start ... Some Mathematical Basics 2
1.1 Review of Simple Series 2
1.2 Review of Double Integrals and Sums 4
1.3 Fundamental Theorem of Calculus 7
1.4 Review of Taylor Series and Other Limits 8
1.5 A Little Combinatorics Ripasso: combinatoria 11
1.6 Review of Asymptotic Notation 15
1.6.1 Big-O and Little-o 15
1.6.2 Big-Omega and Little-omega 17
1.6.3 Big-Theta 18
1.7 Exercises 18

2 Probability on Events Probabilità di base 21
2.1 Sample Space and Events 21
2.2 Probability Defined on Events 22
2.3 Conditional Probabilities on Events 24
2.4 Independent Events 27
2.5 Law of Total Probability 30
2.6 Bayes' Law 32
2.7 Exercises 34

Part II Discrete Random Variables Variabili casuali dicrete

3 Common Discrete Random Variables Alcune distribuzioni discrete 44
3.1 Random Variables 44
3.2 Common Discrete Random Variables 45
3.2.1 The Bernoulli(p) Random Variable 46
3.2.2 The Binomial(n,p) Random Variable 47
3.2.3 The Geometric(p) Random Variable 48

Mor Harchol-Balter. *Introduction to Probability for Computing,*
Cambridge University Press, 2024. Not for distribution.

3.2.4 The Poisson(λ) Random Variable 49
3.3 Multiple Random Variables and Joint Probabilities 50
3.4 Exercises 54

4 Expectation Valore atteso (o media) 58
4.1 Expectation of a Discrete Random Variable 58
4.2 Linearity of Expectation 63
4.3 Conditional Expectation 67
4.4 Computing Expectations via Conditioning 72
4.5 Simpson's Paradox 74
4.6 Exercises 76

5 Variance, Higher Moments, and Random Sums Varianza, altri momenti e somme aleatorie 83
5.1 Higher Moments 83
5.2 Variance 85
5.3 Alternative Definitions of Variance 86
5.4 Properties of Variance 88
5.5 Summary Table for Discrete Distributions 91
5.6 Covariance 91
5.7 Central Moments 92
5.8 Sum of a Random Number of Random Variables 93
5.9 Tails 97
5.9.1 Simple Tail Bounds 98
5.9.2 Stochastic Dominance 99
5.10 Jensen's Inequality 102
5.11 Inspection Paradox 104
5.12 Exercises 107

6 z-Transforms 116
6.1 Motivating Examples 116
6.2 The Transform as an Onion 117
6.3 Creating the Transform: Onion Building 118
6.4 Getting Moments: Onion Peeling 120
6.5 Linearity of Transforms 121
6.6 Conditioning 123
6.7 Using z-Transforms to Solve Recurrence Relations 124
6.8 Exercises 128

Part III Continuous Random Variables

7 Continuous Random Variables: Single Distribution 134
7.1 Probability Density Functions Variabili casuali continue 134
7.2 Common Continuous Distributions 137
7.3 Expectation, Variance, and Higher Moments 141

Mor Harchol-Balter. *Introduction to Probability for Computing,*
Cambridge University Press, 2024. Not for distribution.

7.4 Computing Probabilities by Conditioning on a R.V. 143
7.5 Conditional Expectation and the Conditional Density 146
7.6 Exercises 150

8 Continuous Random Variables: Joint Distributions 153
8.1 Joint Densities Distribuzioni congiunte continue 153
8.2 Probability Involving Multiple Random Variables 156
8.3 Pop Quiz 160
8.4 Conditional Expectation for Multiple Random Variables 161
8.5 Linearity and Other Properties 163
8.6 Exercises 163

9 Normal Distribution La distribuzione normale o Gaussiana 170
9.1 Definition 170
9.2 Linear Transformation Property 172
9.3 The Cumulative Distribution Function 173
9.4 Central Limit Theorem 176
9.5 Exercises 178

10 Heavy Tails: The Distributions of Computing Code pesanti 181
10.1 Tales of Tails 181
10.2 Increasing versus Decreasing Failure Rate 183
10.3 UNIX Process Lifetime Measurements 186
10.4 Properties of the Pareto Distribution 187
10.5 The Bounded-Pareto Distribution 189
10.6 Heavy Tails 189
10.7 The Benefits of Active Process Migration 190
10.8 From the 1990s to the 2020s 191
10.9 Pareto Distributions Are Everywhere 192
10.10 Summary Table for Continuous Distributions 194
10.11 Exercises 194

11 Laplace Transforms 198
11.1 Motivating Example 198
11.2 The Transform as an Onion 198
11.3 Creating the Transform: Onion Building 200
11.4 Getting Moments: Onion Peeling 201
11.5 Linearity of Transforms 203
11.6 Conditioning 203
11.7 Combining Laplace and z-Transforms 204
11.8 One Final Result on Transforms 205
11.9 Exercises 206

Mor Harchol-Balter. *Introduction to Probability for Computing,*
Cambridge University Press, 2024. Not for distribution.

Part IV Computer Systems Modeling and Simulation

12 The Poisson Process 210
12.1 Review of the Exponential Distribution 210
12.2 Relating the Exponential Distribution to the Geometric 211
12.3 More Properties of the Exponential 213
12.4 The Celebrated Poisson Process 216
12.5 Number of Poisson Arrivals during a Random Time 219
12.6 Merging Independent Poisson Processes 220
12.7 Poisson Splitting 221
12.8 Uniformity 224
12.9 Exercises 225

13 Generating Random Variables for Simulation 229
13.1 Inverse Transform Method 229
13.1.1 The Continuous Case 230
13.1.2 The Discrete Case 231
13.2 Accept–Reject Method 232
13.2.1 Discrete Case 233
13.2.2 Continuous Case 234
13.2.3 A Harder Problem 238
13.3 Readings 238
13.4 Exercises 238

14 Event-Driven Simulation 240
14.1 Some Queueing Definitions 240
14.2 How to Run a Simulation 242
14.3 How to Get Performance Metrics from Your Simulation 244
14.4 More Complex Examples 247
14.5 Exercises 249

Part V Statistical Inference

15 Estimators for Mean and Variance 255
15.1 Point Estimation 255
15.2 Sample Mean 256
15.3 Desirable Properties of a Point Estimator 256
15.4 An Estimator for Variance 259
15.4.1 Estimating the Variance when the Mean is Known 259
15.4.2 Estimating the Variance when the Mean is Unknown 259
15.5 Estimators Based on the Sample Mean 261
15.6 Exercises 263
15.7 Acknowledgment 264

Mor Harchol-Balter. *Introduction to Probability for Computing,*
Cambridge University Press, 2024. Not for distribution.

16 Classical Statistical Inference 265
16.1 Towards More General Estimators 265
16.2 Maximum Likelihood Estimation 267
16.3 More Examples of ML Estimators 270
16.4 Log Likelihood 271
16.5 MLE with Data Modeled by Continuous Random Variables 273
16.6 When Estimating More than One Parameter 276
16.7 Linear Regression 277
16.8 Exercises 283
16.9 Acknowledgment 284

17 Bayesian Statistical Inference 285
17.1 A Motivating Example 285
17.2 The MAP Estimator 287
17.3 More Examples of MAP Estimators 290
17.4 Minimum Mean Square Error Estimator 294
17.5 Measuring Accuracy in Bayesian Estimators 299
17.6 Exercises 301
17.7 Acknowledgment 304

Part VI Tail Bounds and Applications

18 Tail Bounds 306
18.1 Markov's Inequality 307
18.2 Chebyshev's Inequality 308
18.3 Chernoff Bound 309
18.4 Chernoff Bound for Poisson Tail 311
18.5 Chernoff Bound for Binomial 312
18.6 Comparing the Different Bounds and Approximations 313
18.7 Proof of Chernoff Bound for Binomial: Theorem 18.4 315
18.8 A (Sometimes) Stronger Chernoff Bound for Binomial 316
18.9 Other Tail Bounds 318
18.10 Appendix: Proof of Lemma 18.5 319
18.11 Exercises 320

19 Applications of Tail Bounds: Confidence Intervals and Balls and Bins 327
19.1 Interval Estimation 327
19.2 Exact Confidence Intervals 328
19.2.1 Using Chernoff Bounds to Get Exact Confidence Intervals 328
19.2.2 Using Chebyshev Bounds to Get Exact Confidence Intervals 331
19.2.3 Using Tail Bounds to Get Exact Confidence Intervals in General Settings 332
19.3 Approximate Confidence Intervals 334

Mor Harchol-Balter. *Introduction to Probability for Computing,*
Cambridge University Press, 2024. Not for distribution.

19.4 Balls and Bins 337
19.5 Remarks on Balls and Bins 341
19.6 Exercises 341

20 Hashing Algorithms 346
20.1 What is Hashing? 346
20.2 Simple Uniform Hashing Assumption 348
20.3 Bucket Hashing with Separate Chaining 349
20.4 Linear Probing and Open Addressing 352
20.5 Cryptographic Signature Hashing 355
20.6 Remarks 360
20.7 Exercises 360

Part VII Randomized Algorithms

21 Las Vegas Randomized Algorithms 364
21.1 Randomized versus Deterministic Algorithms 364
21.2 Las Vegas versus Monte Carlo 366
21.3 Review of Deterministic Quicksort 367
21.4 Randomized Quicksort 368
21.5 Randomized Selection and Median-Finding 370
21.6 Exercises 373

22 Monte Carlo Randomized Algorithms 383
22.1 Randomized Matrix-Multiplication Checking 383
22.2 Randomized Polynomial Checking 387
22.3 Randomized Min-Cut 389
22.4 Related Readings 394
22.5 Exercises 394

23 Primality Testing 403
23.1 Naive Algorithms 403
23.2 Fermat's Little Theorem 404
23.3 Fermat Primality Test 408
23.4 Miller–Rabin Primality Test 410
23.4.1 A New Witness of Compositeness 410
23.4.2 Logic Behind the Miller–Rabin Test 411
23.4.3 Miller–Rabin Primality Test 413
23.5 Readings 415
23.6 Appendix: Proof of Theorem 23.9 415
23.7 Exercises 417

Part VIII Discrete-Time Markov Chains

Mor Harchol-Balter. *Introduction to Probability for Computing,*
Cambridge University Press, 2024. Not for distribution.

24 Discrete-Time Markov Chains: Finite-State 420
24.1 Our First Discrete-Time Markov Chain 420
24.2 Formal Definition of a DTMC 421
24.3 Examples of Finite-State DTMCs 422
24.3.1 Repair Facility Problem 422
24.3.2 Umbrella Problem 423
24.3.3 Program Analysis Problem 424
24.4 Powers of $\mathbf{P}$: n-Step Transition Probabilities 425
24.5 Limiting Probabilities 426
24.6 Stationary Equations 428
24.7 The Stationary Distribution Equals the Limiting Distribution 429
24.8 Examples of Solving Stationary Equations 432
24.9 Exercises 433

25 Ergodicity for Finite-State Discrete-Time Markov Chains 438
25.1 Some Examples on Whether the Limiting Distribution Exists 439
25.2 Aperiodicity 441
25.3 Irreducibility 442
25.4 Aperiodicity plus Irreducibility Implies Limiting Distribution 443
25.5 Mean Time Between Visits to a State 448
25.6 Long-Run Time Averages 450
25.6.1 Strong Law of Large Numbers 452
25.6.2 A Bit of Renewal Theory 454
25.6.3 Equality of the Time Average and Ensemble Average 455
25.7 Summary of Results for Ergodic Finite-State DTMCs 456
25.8 What If My DTMC Is Irreducible but Periodic? 456
25.9 When the DTMC Is Not Irreducible 457
25.10 An Application: PageRank 458
25.10.1 Problems with Real Web Graphs 461
25.10.2 Google's Solution to Dead Ends and Spider Traps 462
25.10.3 Evaluation of the PageRank Algorithm and Practical Considerations 463
25.11 From Stationary Equations to Time-Reversibility Equations 464
25.12 Exercises 469

26 Discrete-Time Markov Chains: Infinite-State 479
26.1 Stationary = Limiting 479
26.2 Solving Stationary Equations in Infinite-State DTMCs 480
26.3 A Harder Example of Solving Stationary Equations in Infinite-State DTMCs 483
26.4 Ergodicity Questions 484
26.5 Recurrent versus Transient: Will the Fish Return to Shore? 487
26.6 Infinite Random Walk Example 490
26.7 Back to the Three Chains and the Ergodicity Question 492

Catene di Markov a tempo discreto

Mor Harchol-Balter. *Introduction to Probability for Computing,* Cambridge University Press, 2024. Not for distribution.

26.7.1 Figure 26.8(a) is Recurrent 492
26.7.2 Figure 26.8(b) is Transient 492
26.7.3 Figure 26.8(c) is Recurrent 494
26.8 Why Recurrence Is Not Enough 494
26.9 Ergodicity for Infinite-State Chains 496
26.10 Exercises 498

27 A Little Bit of Queueing Theory 510
27.1 What Is Queueing Theory? 510
27.2 A Single-Server Queue 511
27.3 Kendall Notation 513
27.4 Common Performance Metrics 514
27.4.1 Immediate Observations about the Single-Server Queue 515
27.5 Another Metric: Throughput 516
27.5.1 Throughput for $M/G/k$ 517
27.5.2 Throughput for Network of Queues with Probabilistic Routing 518
27.5.3 Throughput for Network of Queues with Deterministic Routing 519
27.5.4 Throughput for Finite Buffer 520
27.6 Utilization 520
27.7 Introduction to Little's Law 521
27.8 Intuitions for Little's Law 522
27.9 Statement of Little's Law 524
27.10 Proof of Little's Law 525
27.11 Important Corollaries of Little's Law 527
27.12 Exercises 531

References 539

Index 544

Mor Harchol-Balter. *Introduction to Probability for Computing,*
Cambridge University Press, 2024. Not for distribution.